




Optimizations in preparation stage of IdMapper






 Merged

 tinwelint merged 2 commits into neo4j:2.2 from tinwelint:2.2-parallel-sort-opti Apr 12, 2015

 Conversation 0

 Commits 2

 Files changed 15

+274 -223     



tinwelint commented Apr 9, 2015

Collaborator

A previously costly initialization loop before sorting started now runs in parallel on all available CPUs. Also both collision detection and binarySearch algorithms does fewer data cache comparisons for a bit better performance.

Locally measured to cut the time spent in Prepare node index stage in half.

Labels

2.2

kernel

Milestone

No milestone

Assignee


 digitalstain

2 participants



tinwelint added some commits Apr 1, 2015



 Initialization of tracker cache in ParallelSort runs in parallel ...

0cac57a



 Reduces number of data comparisons in IdMapper

 967c16a



tinwelint added kernel 2.2 labels Apr 9, 2015



digitalstain was assigned by tinwelint Apr 10, 2015



tinwelint merged commit 902d898 into neo4j:2.2 Apr 12, 2015

1 check passed

[View details](#)

[Sign up for free](#)

to join this conversation on GitHub. Already have an account? [Sign in to comment](#)

