



This repository Search

Explore Features Enterprise Blog

Sign up

Sign in



neo4j / neo4j

Watch

238

Star

1,885

Fork

730

# Abnormal performance drop with length increase of "in" operator argument

New issue

Closed

rsk700 opened this issue Apr 13, 2015 · 8 comments



rsk700 commented Apr 13, 2015

Was tested with 2.2, but looks like same valid for 2.1.7

STR:

Test data:

```
unwind range(1,100000) as r
create (:page {
  category: 'category-' + str(r)
})
```

Check performance:

```
match (p:page)
where p.category in ['category-1']
return count(*)
```

58 ms

```
match (p:page)
where p.category in ['category-1', 'cat-1', 'cat-2', 'cat-3', 'cat-4', 'cat-5', 'cat-6', 'cat-7', 'cat-8', 'cat-9', 'cat-10']
return count(*)
```

1059 ms

increasing length of list leads to significant performance drop

Labels

None yet

Milestone

No milestone

Assignee

No one assigned

3 participants



jexp commented Apr 13, 2015

Collaborator

did you create an index on `:page(category)` ?  
otherwise it will have to scan all nodes and for each node do the IN check (which is not optimized using a Set yet)



jexp commented Apr 13, 2015

Collaborator

The second query executes in less than 40ms for me, please also note you misspelt the "category-xxx" values for everything > 1



jexp commented Apr 13, 2015

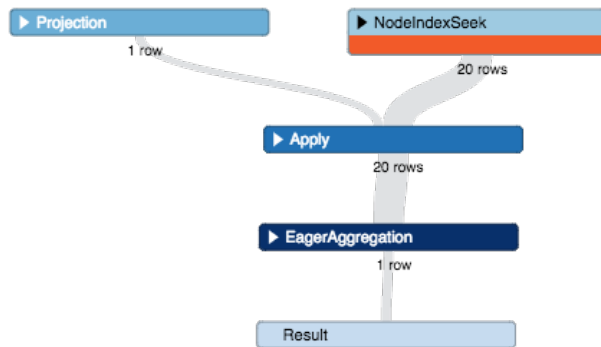
Collaborator

2.2 has an issue where it doesn't automatically use the index when the collection is an expression or parameter, but you can force index usage like this:

```

profile
with [r in range(1,20) | 'category-'+r] as cats
match (p:page)
using index p:page(category)
where p.category in cats
return count(*)

```



**jexp** closed this Apr 13, 2015



**rsk700** commented Apr 13, 2015

did you create an index on :page(category) ?  
otherwise it will have to scan all nodes and for each node do the IN check (which is not optimized using a Set yet)

yes I know, I didn't create index on purpose, and yes both queries scan all nodes, but second query makes it much much slower

The second query executes in less than 40ms for me, please also note you misspelt the "category-xxx" values for everything > 1

I did not misspell, it made on purpose, all queries return same value "1", "cat-x" is fake values to show issue, only one page "category-1" is counted

Can you say which neo4j version you was using? I have this performance issue on 2 servers.



**rsk700** commented Apr 13, 2015

@**jexp** you misunderstood issue please run queries exactly as I provided, there is definitely performance issue in comparing parameter with list



**jexp** commented Apr 13, 2015

Collaborator

Sure, I said so "IN comparison" across all nodes without an index is **not** an optimized operation, and if you do it for 100k nodes it takes a bit (4s on my machine).

And as you can see in the profile it creates quite a bit of db-hits:



rsk700 commented Apr 13, 2015

Thanks for explanation, I see.

Can it be optimization request ;-)? My point here this operation possible to make just once, from what I see neo4j runs query again and again separately for each member of list:

```
2 * 100000 * 21 = 4 200 000 db hits
```

but it can be one operation similar to this:

```
match (p:page)
where p.category =~ 'category-1|cat-1|cat-2|cat-3|cat-4|cat-5|cat-6|cat-7|cat-8|cat-9|cat-10|cat-11|cat-12|cat-13|cat-14|cat-15|cat-16|cat-17|cat-18|cat-19|cat-20'
return count(*)
```

this gives 200 000 db hits



apivaral commented Apr 13, 2015

@jexp any plans to make "IN comparison" an optimised operation?

