

from Neo4j 1.9.9 to Neo4j 2.2.0 exception querying Lucene

8 posts by 3 authors  



Rita

Apr 16



Hi all,

I am passing from Neo4j 1.9.9 to Neo4j 2.2.0, with embedded mode using java. I have inserted the transactions also for read operations, but now when I query my Lucene indexes as this

```
rhits = index.query("cs", "*");  
out.println("#" + rhits.size());  
rhits.close();
```

as you can see I do not have to iterate over the result, I need only the number of results but this new Neo4j version looks like loading all in memory and I get the following error in the first instruction.

```
Exception in thread "main" java.lang.OutOfMemoryError: Java heap space  
    at org.neo4j.collection.primitive.hopscotch.IntArrayBasedKeyTable.initializeTable(IntArrayBasedKeyTable.java:54)  
    at org.neo4j.collection.primitive.hopscotch.IntArrayBasedKeyTable.<init>(IntArrayBasedKeyTable.java:48)  
    at org.neo4j.collection.primitive.hopscotch.LongKeyTable.<init>(LongKeyTable.java:27)  
    at org.neo4j.collection.primitive.Primitive.longSet(Primitive.java:66)  
    at org.neo4j.kernel.impl.coreapi.LegacyIndexProxy$1.<init>(LegacyIndexProxy.java:296)  
    at org.neo4j.kernel.impl.coreapi.LegacyIndexProxy.wrapIndexHits(LegacyIndexProxy.java:294)  
    at org.neo4j.kernel.impl.coreapi.LegacyIndexProxy.query(LegacyIndexProxy.java:352)
```

I never get this with older versions of Neo4j! I always did this operation until version 1.9.9. Could you please help me to avoid this? Is it a bug of library implementation or I have to change the way of querying?

Thanks in advance,
Rita



Michael Hunger

Apr 16



Other recipients: mat...@neotechnology.com

This seems to be a like a bug.

How much heap do you have?

Could you raise an issue on github.com/neo4j/neo4j/issues ?

Thanks so much

Michael

--

You received this message because you are subscribed to the Google Groups "Neo4j" group. To unsubscribe from this group and stop receiving emails from it, send an email to [neo4j+un...@googlegroups.com](mailto:neo4j+unsubscribe@googlegroups.com). For more options, visit <https://groups.google.com/d/optout>.



Rita

Apr 16



Other recipients: mat...@neotechnology.com

Thank you for the reply Michael. I have just published the issue. I was using as usual -Xmx4g. Now I've just tried with 6GB and open and close a single transaction for every query like that on the different indexes and I do not get this exception. So now it has more need of memory for the same operation. I try on other cases. Tell me if there are news on the issue please.
Thank you.

Regards

Rita

- show quoted text -



Rita

Apr 17



Other recipients: mat...@neotechnology.com

I am very sorry to point out that those Lucene queries actually have a changeable behaviour, usually they are slower than the past, and in most case I get that error again (using 6GB heap of 8GB total RAM).

I am updating the graph with delete and update of nodes and relationships. I decreased the number of operations per transaction but most of times I still got this error.

The insertion with Batch Inserter instead seems to be ok!

What do you suggest me please? Keep 1.9.9 version or upgrade to 2.2.1 could be a solution? The new version covers packages related to these problems?

Thanks in advance

Rita

- show quoted text -

Mattias Persson

Apr 20



Other recipients: mat...@neotechnology.com

Hi, the IndexHits instance returned isn't putting everything in that set right away, but over time to avoid returning duplicates when combining transaction state and store state. Looking at it right now I see that this can be made much better by returning hits from transaction state first, putting `_only_` those ids into that set and then comparing - but not adding - when iterating returning ids from store.

Let me see if I can get around fixing that soon...

- show quoted text -



Rita

Apr 20



Other recipients: mat...@neotechnology.com

Hi Mattias,
thank you for the reply. I opened this issue <https://github.com/neo4j/neo4j/issues/4444> . So you confirm that the java heap space error is caused by a possible bug in index handling. In this case I'm going to wait for news about a fix to have a retry in future, I hope soon!
Thank you for your work!

Rita

- show quoted text -



Rita

May 6



Other recipients: mat...@neotechnology.com

Thank you very much for the bug fix <https://github.com/neo4j/neo4j/pull/4509>
When will be possible to have this solution into a new release?

Rita

- show quoted text -

Mattias Persson

May 12



Other recipients: mat...@neotechnology.com

It should be included in 2.2.2 when that gets released, also the next 2.3 milestone (2.3-M02)

- show quoted text -