

Ten. Million. Questions. Let's celebrate [all we've done together](#).

Stack Overflow is a question and answer site for professional and enthusiast programmers. It's 100% free.

Take the 2-minute tour ×

Neo4j Performance Challenge - How to Improve?

▲
3

▼
★
2

I've been wrangling with Neo4J for the last few weeks, trying to resolve some extremely challenging performance problems. At this point, I need some additional help because I can't determine how to move forward.

I have a graph with a total of approx 12.5 Million nodes and 64 Million relationships. The purpose of the graph is going to be analyzing suspicious financial behavior, so it is customers, accounts, transactions, etc.

Here is an example of the performance challenge:

- This query for total nodes takes 96,064ms to complete, which is extremely long.

```
neo4j-sh (?)$ MATCH (n) RETURN count(n);
+-----+
| count(n) |
+-----+
| 12519940 |
+-----+
1 row
96064 ms
```

- The query for total relationships takes 919,449ms to complete, which seems silly.

```
neo4j-sh (?)$ MATCH ()-[r]-() return count(r);
+-----+
| count(r) |
+-----+
| 64062508 |
+-----+
1 row
919449 ms
```

- I have 6.6M Transaction Nodes. When I attempt to search for transactions that have an amount above \$8,000, the query takes 653,637ms also way too long.

```
neo4j-sh (?)$ MATCH (t:Transaction) WHERE t.amount > 8000.00 return count(t);
+-----+
| count(t) |
+-----+
| 10696 |
+-----+
1 row
653637 ms
```

Relevant Schema

```
ON :Transaction(baseamount)  ONLINE
ON :Transaction(type)        ONLINE
ON :Transaction(amount)      ONLINE
ON :Transaction(currency)    ONLINE
ON :Transaction(basecurrency) ONLINE
ON :Transaction(transactionid) ONLINE (for uniqueness constraint)
```

Profile of Query:

```
neo4j-sh (?)$ PROFILE MATCH (t:Transaction) WHERE t.amount > 8000.00 return count(t)
+-----+
| count(t) |
+-----+
| 10696 |
+-----+
```

asked 4 months ago

viewed 96 times

active 4 months ago









Blog

 [Why Stack Overflow is a Good Workplace for Women](#)

Related

- 0 [Improve Neo4j Cypher Performance On Lengthy Match](#)
- 0 [How to improve Neo4j 2.0 cypher/ExecutionResult performance under heavy load?](#)
- 1 [neo4j improving cypher query performance](#)
- 0 [How to return random records in Neo4j using Cypher?](#)
- 0 [how to improve Neo4J performance in creating edges?](#)
- 1 [Neo4j: how does relationships concentration influence performance?](#)
- 1 [neo4j cypher - Unable to execute query \[400\], 'StackOverflowError' exception](#)
- 1 [Neo4j write performance improvement](#)
- 1 [Neo4j - don't know how to improve cypher query](#)
- 1 [Neo4j - how to improve queries response time?](#)

Hot Network Questions

-  [The Programming Language Quiz](#)
-  [Combining several symmetric ciphers using XOR](#)
-  [Does touching the tent outer wall from the inside make it leak?](#)
-  [Why are carpenter's pencils flat?](#)
-  [An "outside the box" sequence](#)
-  [Manager sounds upset every time I inform him of a \(minor\) obstacle](#)
-  [When does the designer pay for a reprint](#)
-  [Lunch with the CEO](#)

```

1 row
ColumnFilter
|
+EagerAggregation
|
+Filter
|
+NodeByLabel

```

Operator	Rows	DbHits	Identifiers	
ColumnFilter	1	0		keep colu
EagerAggregation	1	0		
Filter	10696	13216382		Property(t,amount(62)) > {
NodeByLabel	6608191	6608192	t, t	

- I am running these in the neo4j shell.
- The performance challenges here are starting to create substantial doubt about whether I can even use Neo4J, and seem opposite of the potential the platform offers.
- I am fully admit that I may have misconfigured something (I'm relatively new to Neo4J), so guidance on what to fix or what to look at is much appreciated.

Here are details of my setup:

System: Linux, Ubuntu, 16GB RAM, 3.5 i5 Proc, 256GB SSD HD

CPU

```

$ cat /proc/cpuinfo
processor       : 0
vendor_id     : GenuineIntel
cpu family    : 6
model         : 60
model name    : Intel(R) Core(TM) i5-4690K CPU @ 3.50GHz
stepping      : 3
microcode     : 0x12
cpu MHz       : 4230.625
cache size    : 6144 KB

```

Memory

```

$ cat /proc/meminfo
MemTotal:      16115020 kB
MemFree:       224856 kB
MemAvailable:  8807160 kB
Buffers:       124356 kB
Cached:        8429964 kB
SwapCached:    8388 kB

```

Disk

```

$ df -h
Filesystem      Size  Used Avail Use% Mounted on
/dev/mapper/data1--vg-root 219G  32G  177G  16% /

```

Neo4J.properties

```

neostore.nodestore.db.mapped_memory=200M
neostore.relationshipstore.db.mapped_memory=1G
neostore.relationshipgroupstore.db.mapped_memory=200M
neostore.propertystore.db.mapped_memory=500M
neostore.propertystore.db.strings.mapped_memory=500M
neostore.propertystore.db.arrays.mapped_memory=50M
neostore.propertystore.db.index.keys.mapped_memory=200M
relationship_auto_indexing=true

```

Neo4J-Wrapper.properties

- Did a Kamikaze ever impact the hull of a ship, as opposed to the deck?
- How can I grant a non-Super User group permission to force check-in of articles?
- Which Star Trek character appeared on screen with the most different ranks?
- Funding government by only printing more money
- Why don't decoupling/bypass capacitors need resistors to perform their function, like regular filters?
- Weird performance increase in simple benchmark
- Using Emergency Fund to Sell Upside-down Car
- Sign that word!
- Why did Holly decide to simply wait out the radioactivity?
- Which type of drill bit for which jobs?
- If a stock doesn't pay dividends, then why is the stock worth anything?
- Carrying a handgun in other countries with a US concealed carry permit
- How to show pagination "go to page" text box?
- As a professional photographer, how can one handle a wedding photographer being rude or demanding in a way that will compromise overall image quality?
- Why moving fan seems transparent?
- Aerofoil that gives reasonably good lift for both flow directions: forward and backward

```

wrapper.java.additional=-Dorg.neo4j.server.properties=conf/neo4j-server.properties
wrapper.java.additional=-Djava.util.logging.config.file=conf/logging.properties
wrapper.java.additional=-Dlog4j.configuration=file:conf/log4j.properties

#*****
# JVM Parameters
#*****

wrapper.java.additional=-XX:+UseConcMarkSweepGC
wrapper.java.additional=-XX:+CMSClassUnloadingEnabled
wrapper.java.additional=-XX:-OmitStackTraceInFastThrow

# Uncomment the following lines to enable garbage collection logging
wrapper.java.additional=-Xloggc:data/log/neo4j-gc.log
wrapper.java.additional=-XX:+PrintGCDetails
wrapper.java.additional=-XX:+PrintGCDateStamps
wrapper.java.additional=-XX:+PrintGCApplicationStoppedTime
wrapper.java.additional=-XX:+PrintPromotionFailure
wrapper.java.additional=-XX:+PrintTenuringDistribution

# Java Heap Size: by default the Java heap size is dynamically
# calculated based on available system resources.
# Uncomment these lines to set specific initial and maximum
# heap size in MB.
wrapper.java.initmemory=4096
wrapper.java.maxmemory=6144

```

Other:

- Changed the open file settings for Linux to 40k
- I am not running anything else on this machine, no X Windows, no other DB server. Here is a snippet of top while running a query:

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
15785	neo4j		20	0 12.192g	8.964g	2.475g	S	100.2	58.3	227:50.98	java
1	root	20	0	33464	2132	1140	S	0.0	0.0	0:02.36	init
2	root	20	0	0	0	0	S	0.0	0.0	0:00.01	kthreadd

- The total file size in the graph.db directory is:

```

data/graph.db$ du --max-depth=1 -h
1.9G    ./schema
36K     ./index
26G     .

```

- Data loading was extremely hit or miss. Some merges would take less than 60 seconds (Even for ~200 to 300K inserts), while some merges would last for over 3 hours (11,898,514ms for a CSV file with 189,999 rows merging on one date)
- I get constant GC thread blocking:

```

2015-03-27 14:56:26.347+0000 WARN [o.n.k.EmbeddedGraphDatabase]: GC Monitor: Ap
2015-03-27 14:56:39.011+0000 WARN [o.n.k.EmbeddedGraphDatabase]: GC Monitor: Ap
2015-03-27 14:56:57.533+0000 WARN [o.n.k.EmbeddedGraphDatabase]: GC Monitor: Ap
2015-03-27 14:57:17.345+0000 WARN [o.n.k.EmbeddedGraphDatabase]: GC Monitor: Ap
2015-03-27 14:57:29.955+0000 WARN [o.n.k.EmbeddedGraphDatabase]: GC Monitor: Ap
2015-03-27 14:58:14.311+0000 WARN [o.n.k.EmbeddedGraphDatabase]: GC Monitor: Ap

```

Please let me know if I should add anything else that would be salient to the discussion

Update 1

Thank you very much for your help, I just moved so I was delayed in responding.

1. Size of Neostore Files:

```

/data/graph.db$ ls -lah neostore.*
-rw-rw-r-- 1 neo4j neo4j 9 Apr 2 13:03 neostore.id
-rw-rw-r-- 1 neo4j neo4j 110 Apr 2 13:03 neostore.labeltokenstore.db
-rw-rw-r-- 1 neo4j neo4j 9 Apr 2 13:03 neostore.labeltokenstore.db.id
-rw-rw-r-- 1 neo4j neo4j 874 Apr 2 13:03 neostore.labeltokenstore.db.names
-rw-rw-r-- 1 neo4j neo4j 9 Apr 2 13:03 neostore.labeltokenstore.db.names.id
-rw-rw-r-- 1 neo4j neo4j 200M Apr 2 13:03 neostore.nodestore.db
-rw-rw-r-- 1 neo4j neo4j 41 Apr 2 13:03 neostore.nodestore.db.id
-rw-rw-r-- 1 neo4j neo4j 68 Apr 2 13:03 neostore.nodestore.db.labels
-rw-rw-r-- 1 neo4j neo4j 9 Apr 2 13:03 neostore.nodestore.db.labels.id
-rw-rw-r-- 1 neo4j neo4j 2.8G Apr 2 13:03 neostore.propertystore.db
-rw-rw-r-- 1 neo4j neo4j 128 Apr 2 13:03 neostore.propertystore.db.arrays
-rw-rw-r-- 1 neo4j neo4j 9 Apr 2 13:03 neostore.propertystore.db.arrays.id
-rw-rw-r-- 1 neo4j neo4j 9 Apr 2 13:03 neostore.propertystore.db.id
-rw-rw-r-- 1 neo4j neo4j 720 Apr 2 13:03 neostore.propertystore.db.index
-rw-rw-r-- 1 neo4j neo4j 9 Apr 2 13:03 neostore.propertystore.db.index.id

```

```

-rw-rw-r-- 1 neo4j neo4j 3.1K Apr 2 13:03 neostore.propertystore.db.index.keys
-rw-rw-r-- 1 neo4j neo4j 9 Apr 2 13:03 neostore.propertystore.db.index.keys.
-rw-rw-r-- 1 neo4j neo4j 1.7K Apr 2 13:03 neostore.propertystore.db.strings
-rw-rw-r-- 1 neo4j neo4j 9 Apr 2 13:03 neostore.propertystore.db.strings.id
-rw-rw-r-- 1 neo4j neo4j 47M Apr 2 13:03 neostore.relationshipgroupstore.db
-rw-rw-r-- 1 neo4j neo4j 9 Apr 2 13:03 neostore.relationshipgroupstore.db.id
-rw-rw-r-- 1 neo4j neo4j 1.1G Apr 2 13:03 neostore.relationshipstore.db
-rw-rw-r-- 1 neo4j neo4j 1.6M Apr 2 13:03 neostore.relationshipstore.db.id
-rw-rw-r-- 1 neo4j neo4j 165 Apr 2 13:03 neostore.relationshiptypestore.db
-rw-rw-r-- 1 neo4j neo4j 9 Apr 2 13:03 neostore.relationshiptypestore.db.id
-rw-rw-r-- 1 neo4j neo4j 1.3K Apr 2 13:03 neostore.relationshiptypestore.db.nam
-rw-rw-r-- 1 neo4j neo4j 9 Apr 2 13:03 neostore.relationshiptypestore.db.nam
-rw-rw-r-- 1 neo4j neo4j 3.5K Apr 2 13:03 neostore.schemastore.db
-rw-rw-r-- 1 neo4j neo4j 25 Apr 2 13:03 neostore.schemastore.db.id

```

- I read that mapped memory settings are replaced by another cache, and I have commented out those settings.
- Java Profiler

```

Jvmtop 0.8.0 alpha - 16:12:59, amd64, 4 cpus, Linux 3.16.0-33, load avg 0.3
http://code.google.com/p/jvmtop

```

Profiling PID 4260: org.neo4j.server.Bootstrapper

```

68.67% ( 14.01s) org.neo4j.kernel.impl.nioneo.store.StoreFileChannel.read
18.73% ( 3.82s) org.neo4j.kernel.impl.nioneo.store.StoreFailureException
2.86% ( 0.58s) org.neo4j.kernel.impl.cache.ReferenceCache.put()
1.11% ( 0.23s) org.neo4j.helpers.Counter.inc()
0.87% ( 0.18s) org.neo4j.kernel.impl.cache.ReferenceCache.get()
0.65% ( 0.13s) org.neo4j.cypher.internal.compiler.v2_1.parser.Literals$
0.63% ( 0.13s) org.parboiled.scala.package$.getCurrentRuleMethod()
0.62% ( 0.13s) scala.collection.mutable.OpenHashMap.<init>()
0.62% ( 0.13s) scala.collection.mutable.AbstractSeq.<init>()
0.62% ( 0.13s) org.neo4j.kernel.impl.cache.AutoLoadingCache.get()
0.61% ( 0.13s) scala.collection.TraversableLike$$anonfun$map$1.apply()
0.61% ( 0.12s) org.neo4j.kernel.impl.transaction.TxManager.assertTmOk()
0.61% ( 0.12s) org.neo4j.cypher.internal.compiler.v2_1.commands.EntityP
0.61% ( 0.12s) scala.collection.AbstractTraversable.<init>()
0.61% ( 0.12s) scala.collection.immutable.List.toString()
0.60% ( 0.12s) org.neo4j.kernel.impl.nioneo.store.NodeStore.getRecord()
0.57% ( 0.12s) org.neo4j.kernel.impl.transaction.TxManager.getTransacti
0.37% ( 0.08s) org.parboiled.scala.Parser$class.rule()
0.06% ( 0.01s) scala.util.DynamicVariable.value()

```

neo4j cypher

share improve this question

edited Apr 2 at 20:15

asked Mar 27 at 15:03

 BC Smith
16 • 3

can you show the individual file sizes for `neostore.*` – Michael Hunger Mar 28 at 21:11

all your queries are full scans and are so mostly bound by your disk speed, is each of those the first query you run? Can you test the disk performance? And the disk scheduler which should be noop or deadline. – Michael Hunger Mar 28 at 21:13

The garbage collection also doesn't look good. Can you take a thread dump (kill -3 <pid>) and connect a profiler to your neo instance and report back e.g. `jvmtop.sh --profile <pid>` github.com/mneedham/jvmtop/tree/master/jvmtop – Michael Hunger Mar 28 at 21:14

Also your merge and create performance doesn't sound right. There seems to be something really off with your machine. – Michael Hunger Mar 28 at 21:15

but in general full scans is not what a graph database is built for and according to your description also not what your use-case is about. – Michael Hunger Mar 28 at 21:16

add a comment

1 Answer

active

oldest

votes



2

Unfortunately the schema indexes (aka those created using `CREATE INDEX ON :Label(property)`) do not yet support larger than/smaller than conditions. Therefore Neo4j falls back to scan all nodes with the given label and filter on their properties. This is of course expensive.

I do see two different approaches to tackle this:

- 1) If your condition does always have a pre-defined maximum granularity e.g. 10s of USDs, you can build up an "amount-tree" similar to a time-tree (see <http://graphaware.com/neo4j/2014/08/20/graphaware-neo4j-timetree.html>).
- 2) if you don't know the granularity upfront the other option is to setup a manual or auto index for the amount property, see <http://neo4j.com/docs/stable/indexing.html>. The most easy thing is probably using auto index. In `neo4j.properties` set the following options:

```
node_auto_indexing=true
node_keys_indexable=amount
```

Note that this will not automatically add all existing transaction into that index, it just puts those in the index that have been written to since auto indexing is enabled.

You can do a explicit range query on the auto index using

```
MATCH t=node:node_auto_index("amount: [6000 TO 99999999]")
RETURN count(t)
```

[share](#) [improve this answer](#)

[add a comment](#)

answered Mar 28 at 11:57



Stefan Armbruster

23k ● 2 ● 28 ● 51

Your Answer

B *I*

Sign up or [log in](#)

Post as a guest

Sign up using Google

Sign up using Facebook

Sign up using Stack Exchange

Name

Email

[Post Your Answer](#)

By posting your answer, you agree to the [privacy policy](#) and [terms of service](#).

Not the answer you're looking for? Browse other questions tagged [neo4j](#) [cypher](#) or [ask your own question](#).

[question feed](#)

TECHNOLOGY			LIFE / ARTS	CULTURE / RECREATION	SCIENCE	OTHER
Stack Overflow	Programmers	Database Administrators	Photography	English Language & Usage	Mathematics	Stack Apps
Server Fault	Unix & Linux	Drupal Answers	Science Fiction & Fantasy	Skeptics	Cross Validated (stats)	Meta Stack Exchange
Super User	Ask Different (Apple)	SharePoint	Graphic Design	Mi Yodeya (Judaism)	Theoretical Computer Science	Area 51
Web Applications	WordPress Development	User Experience	Movies & TV	Travel	Physics	Stack Overflow Careers
Ask Ubuntu	Geographic Information Systems	Mathematica	Seasoned Advice (cooking)	Christianity	MathOverflow	
Webmasters		Salesforce	Home Improvement	Arqade (gaming)	Chemistry	
Game Development	Electrical Engineering	ExpressionEngine@ Answers	Personal Finance & Money	Bicycles	Biology	
TeX - LaTeX	Android Enthusiasts		Academia	Role-playing Games	more (5)	
	Information Security	more (13)	more (9)	more (21)		