

Ten. Million. Questions. Let's celebrate [all we've done together](#).

Stack Overflow is a question and answer site for professional and enthusiast programmers. It's 100% free.

Take the 2-minute tour

## What's the optimum configuration for sub second query latency in Neo4j?

Work on work you love. From home.



stackoverflowcareers

asked 4 months ago

viewed 54 times

active 4 months ago



1

I have imported nodes.tsv ( 350MB, 18M rows, 3 cols) and rels.tsv ( 5GB, 150M rows, 2 cols) using the batch-importer script. These are my batch.properties file entries

- neostore.nodestore.db.mapped\_memory=250M •
- neostore.relationshipstore.db.mapped\_memory=1000M •
- neostore.relationshipgroupstore.db.mapped\_memory=10M •
- neostore.propertystore.db.mapped\_memory=500M •
- neostore.propertystore.db.strings.mapped\_memory=500M •
- neostore.propertystore.db.arrays.mapped\_memory=215M • dump\_configuration=true

I have turned on auto update and auto indexing in neo4j.properties as follows

- allow\_store\_upgrade=true • node\_auto\_indexing=true • node\_keys\_indexable=name,title •
- relationship\_auto\_indexing=true • relationship\_keys\_indexable=sent\_date,has\_read

I'm using neo4j 2.2 version on 64 bit windows server that has 1 TB SSD and 256GB ram. What's the configuration for batch importer and neo4j server that I should use for maximum query and data loading peformance? This query for ex: is timing out in the browser

```
MATCH ()-[r:BELONGS_TO]->() RETURN r
```

neo4j

share improve this question

add a comment

asked Apr 7 at 2:13

user330612

320 1 6 21

Blog

Why Stack Overflow is a Good Workplace for Women

### Related

- 1 Neo4j Configuration
- 4 All queries are slow with neo4j
- 2 Neo4j Browser – query lost on refresh
- 0 Neo4j cache configurations
- 1 Process Crashes when importing 5 million rows to Neo4J
- 2 Neo4j ServerPlugin Configuration
- 1 Sub-graph Id a good candidate for a label in neo4j
- 2 Parametrized like query in neo4j
- 0 Neo4j crashes all the time randomly
- 0 query on Neo4j in java applications

### Hot Network Questions

- Funding government by only printing more money
- Make subscript size smaller (always)
- Prove the theorem on analytic geometry in the picture.
- Carrying a handgun in other countries with a US concealed carry permit
- Capture output without redirection and leave it on terminal too
- Name for a device purposefully put together from faulty parts
- Aerofoil that gives reasonably good lift for both flow directions: forward and backward
- Construct a polynomial with given roots
- Dullness vs. going overboard: Should I be calling people 'enfants terribles' in an academic paper?

### 2 Answers

active oldest votes



1

With millions of nodes that query might be slow no matter what you do, though with the amount of memory you have available maybe it wouldn't be a big deal. This is a good guide for calculating memory settings:

<http://neo4j.com/developer/guide-performance-tuning/>



While you're playing, I would set the query timeout on the server so that your queries can't jam up the server and force you to need to restart it:

<http://neo4j.com/docs/stable/server-configuration.html>

You might try starting with `LIMIT` clauses on your queries so that you can get an idea for how the performance degrades as the `LIMIT` increases.

If you can possibly find a way to limit your query based on node selects that would also be helpful, especially if you can do it by a label or by a label/property combination (which you can index).

Lastly, I would try using `EXPLAIN` in the web console to get an idea for how your queries will be executed:

<http://neo4j.com/docs/2.2.0/how-do-i-profile-a-query.html>

Also you can use `PROFILE`, though that will run the query, so you'll need to be a bit more careful there. You can probably use the `LIMIT` here too to play and see how things work

share improve this answer

answered Apr 7 at 13:39




**Brian Underwood**  
4,364 ● 4 ● 18

Thanks Brian for the info. My next step is to profile the queries. What's strange is that my overall database folder size is 11.6GB after insertion and index creation. So all of this can easily reside in the RAM I have. And so I'm surprised that queries are timing out for a dB that can reside entirely in memory. I want to know if there are any props I should increase in batch.properties or neo4j – [user330612](#) Apr 7 at 15:02

I think the main problem is that queries need to anchor on nodes, and if you're just querying on `()` then it essentially has to search the entire database and check every relationship on every node for `BELONGS_TO` relationships – [Brian Underwood](#) Apr 7 at 16:11

[add a comment](#)



If you have that much RAM:

2

Your memory mapping config is wrong for 2.2



use only this setting: ``dbms.pagecache.memory=20G``

and then provide neo4j with 24G heap in neo4j-wrapper.conf

Use Neo4j enterprise which scales much better.

Disable the auto-indexes they are not used for what you're doing.

Your query doesn't make sense for any use-case:

```
MATCH ()-[r:BELONGS_TO]->() RETURN r
```

Sub-second graph queries always start at a set of concrete start-points (retrieved with an index lookup) and then traverse out from those starting points.

Global scan queries like your's will just pull all data in memory and inefficiently work over it.

Esp. if you return so much data you can't assume sub-second performance. The data volume alone will kill it.

So figure out a label + property-values that you want to start from and then write the query that traverses out from those start points.

If you want to have sub-second on something you're doing you have to go down to the Java API and aggregate there, e.g. with a server extension:

```
int counter=0;
for (Relationship r : GlobalGraphOperations.at(db)) {
    if (r.hasType(Types.BELONGS_TO)) counter++;
}
return counter;
```

[share](#) [improve this answer](#)

answered Apr 7 at 22:22



**Michael Hunger**  
25.2k ● 2 ● 21 ● 41

Thanks Michael. I can change the setting of page cage memory in neo4j.properties files. Where is the .conf file located that I can modify the heap size of vm? Couldn't find any .conf file in my db folder – [user330612](#) Apr 7 at 22:35

it's in the conf folder, if you are on windows check the "Options" button of the starter-app. – [Michael Hunger](#) Apr 9 at 21:01




[add a comment](#)

- Suggested alternatives for that horrible new noun 'nice-to-have'?
- Does touching the tent outer wall from the inside make it leak?
- The Programming Language Quiz
- Can a plane fly backward if its prop's in reverse?
- How can I grant a non-Super User group permission to force check-in of articles?
- Proving that all integers are even or odd
- Seven overlapping circles
- Has any Star Trek character appeared in all three division uniform colors?
- How do I remove Candy Crush Saga from Windows 10?
- How to deal with people looking at their mobile phone during my presentation?
- Can a Glaive be stowed in a Ruby Scabbard?
- How does java.util.EnumSet<E> work?
- Is "user-select: none" bad UX?
- Why are my tomatoes cracking and what can I do?
- What does OWA stand for?

Your Answer

Sign up or [log in](#)

Post as a guest

-  Sign up using Google
-  Sign up using Facebook
-  Sign up using Stack Exchange

<b>Name</b>	<input type="text"/>
<b>Email</b>	<input type="text"/>

Post Your Answer

By posting your answer, you agree to the [privacy policy](#) and [terms of service](#).

Not the answer you're looking for? Browse other questions tagged [neo4j](#) or [ask your own question](#).

 question feed

TECHNOLOGY			LIFE / ARTS	CULTURE / RECREATION	SCIENCE	OTHER
Stack Overflow	Programmers	Database Administrators	Photography	English Language & Usage	Mathematics	Stack Apps
Server Fault	Unix & Linux	Drupal Answers	Science Fiction & Fantasy	Skeptics	Cross Validated (stats)	Meta Stack Exchange
Super User	Ask Different (Apple)	SharePoint	Graphic Design	Mi Yodeya (Judaism)	Theoretical Computer Science	Area 51
Web Applications	WordPress Development	User Experience	Movies & TV	Travel	Physics	Stack Overflow Careers
Ask Ubuntu	Geographic Information Systems	Mathematica	Seasoned Advice (cooking)	Christianity	MathOverflow	
Webmasters		Salesforce		Arcade (gaming)	Chemistry	
Game Development	Electrical Engineering	ExpressionEngine® Answers	Home Improvement	Bicycles	Biology	
TeX - LaTeX	Android Enthusiasts		Personal Finance & Money	Role-playing Games	more (5)	
	Information Security	more (13)	Academia	more (21)		
			more (9)			