Cypher query optimisation - Utilising known properties of nodes

Setup:

Neo4j and Cypher version 2.2.0. I'm querying Neo4j as an in-memory instance in Eclipse created TestGraphDatabaseFactory().newImpermanentDatabase();. I'm using this approach as it seems faster than the embedded version and I assume it has the same functionality. My graph database is randomly generated programmatically with varying numbers of nodes.

Background:

I generate cypher queries automatically. These queries are used to try and identify a single 'target' node. I can limit the possible matches of the queries by using known 'node' properties. I only use a 'name' property in this case. If there is a known name for a node, I can use it to find the node id and use this in the start clause. As well as known names, I also know (for some nodes) if there are names known not to belong to a node. I specify this in the where clause.

The sorts of queries that I am running look like this...

```
START

nvari = node(5)

MATCH

 (target:C5)-[:IN_LOCATION]->(nvara:LOCATION),
(nvara:LOCATION)-[:CONNECTED]->(nvarb:LOCATION),
(nvara:LOCATION)-[:CONNECTED]->(nvarc:LOCATION),
(nvard:LOCATION)-[:CONNECTED]->(nvarc:LOCATION),
(nvard:LOCATION)-[:CONNECTED]->(nvare:LOCATION),
(nvare:LOCATION)-[:CONNECTED]->(nvarf:LOCATION),
(nvarg:LOCATION)-[:CONNECTED]->(nvarf:LOCATION),
(nvarg:LOCATION)-[:CONNECTED]->(nvarh:LOCATION),
(nvari:C4)-[:IN_LOCATION]->(nvarg:LOCATION),
(nvarj:C2)-[:IN_LOCATION]->(nvarg:LOCATION),
(nvare:LOCATION)-[:CONNECTED]->(nvark:LOCATION),
(nvarm:C3)-[:IN_LOCATION]->(nvarg:LOCATION),

WHERE

NOT(nvarj.Name IN ['nf']) AND NOT(nvarm.Name IN ['nb','nj'])

RETURN DISTINCT target
```

Another way to think about this (if it helps), is that this is an isomorphism testing problem where we have some information about how nodes in a query and target graph correspond to each other based on restrictions on labels.

asked 4 months ago

viewed 41 times

active 4 months ago

Question:

With regards to optimisation:

1. Would it help to include relation variables in the match clause? I took them out because the node variables are sufficient to distinguish between relationships but this might slow it down?
2. Should I restructure the match clause to have match/where couples including the where clauses from my previous example first? My expectation is that they can limit possible bindings early on. For example...

   START

   nvari = node(5)

   MATCH

   (nvarj:C2)-[:IN_LOCATION]->(nvarg:LOCATION)

   WHERE NOT(nvarj.Name IN ['nf'])

   MATCH

   (nvarm:C3)-[:IN_LOCATION]->(nvarg:LOCATION)

   WHERE NOT(nvarm.Name IN ['nb','nj'])

   MATCH

   (target:C5)-[:IN_LOCATION]->(nvara:LOCATION), (nvara:LOCATION)-[:CONNECTED]->(nvarb:LOCATION), (nvara:LOCATION)-[:CONNECTED]->(nvarc:LOCATION),

   (nvard:LOCATION)-[:CONNECTED]->(nvarc:LOCATION), (nvard:LOCATION)-[:CONNECTED]->(nvare:LOCATION), (nvare:LOCATION)-[:CONNECTED]->(nvarf:LOCATION), (nvarg:LOCATION)-[:CONNECTED]->(nvarf:LOCATION), (nvarg:LOCATION)-[:CONNECTED]->(nvarh:LOCATION), (nvare:LOCATION)-[:CONNECTED]->(nvark:LOCATION)

   RETURN DISTINCT target

On the side:

3. (Less important but still an interest) If I make each relationship in a match clause an optional match except for relationships containing the target node, would cypher essentially be finding a maximum common sub-graph between the query and the graph data base with the constraint that the MCS contains the target node?

Thanks a lot in advance! I hope I have made my requirements clear but I appreciate that this is not a typical use-case for Neo4j.

neo4j cypher

share improve this question

asked Apr 11 at 15:43
Michael Anslow
69 • 9

add a comment

1 Answer

active oldest votes

1

1. I think querying with node properties is almost always preferable to using relationship properties (if you had a choice), as that opens up the possibility that indexing can help speed up the query.

   As an aside, I would avoid using the `IN` operator if the collection of possible values only has a single element. For example, this snippet: `NOT(nvarj.Name IN ['nf'])`, should be `(nvarj.Name <> 'nf')`. The current versions of Cypher might not use an index for the `IN` operator.

2. Restructuring a query to eliminate undesirable bindings earlier is exactly what you*should* be doing.

3. First of all, you would need to keep using `MATCH` for at least the first relationship in your query (which binds `target`), or else your result would contain a lot of `null` rows -- not very useful.

   But, thinking clearly about this, if all the other relationships were placed in separate `OPTIONAl MATCH` clauses, you'd be essentially saying that you want a match even if*none* of the optional matches succeeded. Therefore, the logical equivalent would be:

   ```
   MATCH (target:C5)-[:IN_LOCATION]->(nvara:LOCATION)
   RETURN DISTINCT target
   ```

   I don't think this is a useful result.

share improve this answer

share improve this answer

answered Apr 12 at 3:46

cybersam

9,941 • 2 • 7 • 21

1. In that case, would chaining multiple inequalities with <> bet better than using NOT with IN? E.g. (nvarj.Name <> 'nf') AND (nvarj.Name <> 'ng') 2. Great, good to know it matches my intuition. I assume this is done in order from the first match declared to the last? 3. Good point. I suppose what I really want is possible valuations for target ordered by how many of the optional match clauses were bound. I suppose I would have to use variables for relationships, return these relationships and create this order manually? – Michael Anslow Apr 22 at 12:09

add a comment

Your Answer

B I

Sign up or log in

Sign up using Google

Sign up using Facebook

Sign up using Stack Exchange

Post as a guest

Name

Email

required, but never shown

Post Your Answer

By posting your answer, you agree to the privacy policy and terms of service.

Not the answer you're looking for? Browse other questions tagged neo4j cypher or ask your own question.

question feed