

Ten. Million. Questions. Let's celebrate [all we've done together](#).

Stack Overflow is a question and answer site for professional and enthusiast programmers. It's 100% free.

Take the 2-minute tour ×

Graph database performance



asked 4 months ago

viewed 65 times

active 4 months ago

Blog

 [Why Stack Overflow is a Good Workplace for Women](#)



1



1

I was reading a book recommended on Neo4j site: <http://neo4j.com/books/graph-databases/> about graph database performance and it said:

"In contrast to relational databases, where join-intensive query performance deteriorates as the dataset gets bigger, with a graph database performance tends to remain relatively constant, even as the dataset grows. This is because queries are localized to a **portion of the graph**. **As a result, the execution time for each query is proportional only to the size of the part of the graph traversed to satisfy that query, rather than the size of the overall graph.**"

So e.g. I want to return only nodes with a label "Doctor, that's localized to a portion of a graph. But my question is how does the database itself know where those nodes are ? In other words, does it not need to traverse all nodes to find out whether or not they satisfy the query and make decision based on that ?

[neo4j](#) [graph-databases](#)

share improve this question

add a comment

asked Apr 14 at 14:22



Frugo

229 ● 4 ● 20

2 Answers

active

oldest

votes



1



Neo4j has a special indexing for node labels so that it can find all nodes for a label without searching all nodes. Beyond that you can:

- Create your own indexes based on node properties (either schema indexes or legacy indexes) in order to find nodes as starting points
- Query by node IDs to find a starting point (though I'd suggest using your own property with an index if you need to identify nodes more permanently)

share improve this answer

add a comment

answered Apr 14 at 15:36



Brian Underwood

4,364 ● 4 ● 18

Work on work you love. From home.



 **stackoverflow**careers



1



In general localized searches mean: **you start from a smallish set of starting points** which can be people, products, places, orders etc.

A portion of the graph that is annotated with a label, often doesn't fall into that category, i.e. **all doctors** are not a smallish set of starting points.

Your query would probably touch a large portion of the graph if you traverse out from *all* doctors to their neighborhoods.

A query like this would be a graph local one:

```
MATCH (:City)-[:name="FRUGO"]->[:DESCRIBE]->(:Doctor)-[:name="DESCRIBE"]->(:Medicine)
```



 **Love this site?**

Get the **weekly newsletter!**

- Top questions and answers
- Important announcements
- Unanswered questions

[Sign up for the newsletter](#)

see an [example newsletter](#)

Related

- 101 [Has anyone used Graph-based Databases \(http://neo4j.org/\)?](#)
- 6 [Preserving nodes and relationships history in a graph database](#)
- 0 [Importing the datasets from the Book "Graph databases"](#)
- 2 [Change management for graph databases?](#)
- 1 [Optimize Neo4j Cypher path finding with limited paths in an undirected graph](#)
- 4 [Efficient way to manipulate a large \(doesn't fit into memory\) graph](#)

```
MATCH (:City {name:"SFO"})<-[[:RESIDES_IN]]-(d:DOCTOR)-[presc:PRESCRIBES]->(m:MEDICINE
RETURN d.name, m.name, sum(presc.amount) as amount
```

share improve this answer

answered Apr 14 at 22:08



Michael Hunger

25.2k ● 2 ● 21 ● 41

add a comment

Your Answer

B *I*

Sign up or [log in](#)

Sign up using Google

Sign up using Facebook

Sign up using Stack Exchange

Post as a guest

Name

Email

required, but never shown

Post Your Answer

By posting your answer, you agree to the [privacy policy](#) and [terms of service](#).

Not the answer you're looking for? Browse other questions tagged [neo4j](#) [graph-databases](#) or [ask your own question](#).

3 Algorithms to find nearest nodes in a graph

0 Graph database modelling: Should i use a collection node to avoid to many rel on a node

0 Neo4j Graph database representation library

2 Is this concept applicable for a graph database?

Hot Network Questions

Why does Scala allow nested data structures like List or Array

Should you always minimize cognitive load

An "outside the box" sequence

Two definitions of Green's function

Weird performance increase in simple benchmark

Password strength of random words and brainflayer

¿Puede "hopefully" traducirse como "ojalá"?

[more hot questions](#)

[question feed](#)

[tour](#) [help](#) [blog](#) [chat](#) [data](#) [legal](#) [privacy policy](#) [work here](#) [advertising info](#) [mobile](#) [contact us](#) [feedback](#)

TECHNOLOGY

[Stack Overflow](#)

[Server Fault](#)

[Super User](#)

[Web Applications](#)

[Ask Ubuntu](#)

[Webmasters](#)

[Game Development](#)

[TeX - LaTeX](#)

[Programmers](#)

[Unix & Linux](#)

[Ask Different \(Apple\)](#)

[WordPress Development](#)

[Geographic Information Systems](#)

[Electrical Engineering](#)

[Android Enthusiasts](#)

[Information Security](#)

[Database Administrators](#)

[Drupal Answers](#)

[SharePoint](#)

[User Experience](#)

[Mathematica](#)

[Salesforce](#)

[ExpressionEngine® Answers](#)

[more \(13\)](#)

LIFE / ARTS

[Photography](#)

[Science Fiction & Fantasy](#)

[Graphic Design](#)

[Movies & TV](#)

[Seasoned Advice \(cooking\)](#)

[Home Improvement](#)

[Personal Finance & Money](#)

[Academia](#)

[more \(9\)](#)

CULTURE / RECREATION

[English Language & Usage](#)

[Skeptics](#)

[Mi Yodeya \(Judaism\)](#)

[Travel](#)

[Christianity](#)

[Arqade \(gaming\)](#)

[Bicycles](#)

[Role-playing Games](#)

[more \(21\)](#)

SCIENCE

[Mathematics](#)

[Cross Validated \(stats\)](#)

[Theoretical Computer Science](#)

[Physics](#)

[MathOverflow](#)

[Chemistry](#)

[Biology](#)

[more \(5\)](#)

OTHER

[Stack Apps](#)

[Meta Stack Exchange](#)

[Area 51](#)

[Stack Overflow Careers](#)

